

# Package ‘fonfDNN’

June 24, 2025

**Type** Package

**Title** Hybrid deep tensor network for function-on-function regression with mixed predictors

**Version** 1.0

**Depends** R (>= 3.5.0), fda, keras, tensorflow, fda.usc, plot3D

**Imports** goffda

**LazyLoad** yes

**ByteCompile** TRUE

**Encoding** UTF-8

**Maintainer** Ufuk Beyaztas <ufukbeyaztas@gmail.com>

**Description** Functions for implementing deep tensor network for function-on-function regression with mixed predictors.

**License** GPL-3

## Contents

air_data . . . . .	1
dgp_mixed . . . . .	3
fonf_fit . . . . .	5
fonf_fnc . . . . .	8
fonf_param_grid . . . . .	10
fonf_predict . . . . .	11
fonf_tune . . . . .	12
<b>Index</b>	<b>15</b>

---

air_data	<i>Air Quality Dataset from Lombardy, Italy (2023–2024)</i>
----------	---

---

## Description

A real-world dataset containing daily air pollution measurements from the Lombardy region in Italy, covering the years 2023 and 2024. This dataset includes functional trajectories of various pollutants recorded over time across 1481 monitoring stations, and serves as the empirical application for the proposed hybrid deep tensor network (HDTN) model for function-on-function regression with mixed predictors.

**Usage**

```
data(air_data)
```

**Format**

A list with the following components:

- y A matrix of dimension  $n \times M$  representing the functional response curves (daily mean  $PM_{2.5}$  concentrations), where  $n$  is the number of stations and  $M$  is the number of grid points.
- x A list of length 5, where each element is an  $n \times M$  matrix corresponding to a functional predictor:
  1. Daily mean  $NO_2$  concentration,
  2. Daily maximum 1-hour  $NO_2$  concentration,
  3. Daily maximum 8-hour  $O_3$  concentration,
  4. Daily maximum 1-hour  $O_3$  concentration,
  5. Daily mean  $PM_{10}$  concentration.
- xsc1 A matrix of dimension  $n \times 3$  containing scalar predictors:

**Details**

This dataset is sourced from the ARPA Lombardy regional environmental monitoring agency and was accessed using the ARPALData package. It is used in the empirical section of the paper to demonstrate the effectiveness of the HDTN model in predicting fine particulate matter concentrations based on mixed-type predictors (functional and scalar).

**Value**

A list containing the functional response and covariate data required to fit the HDTN model, including:

y	Functional response ( $PM_{2.5}$ curves)
x	List of 5 functional covariate matrices
xsc1	Matrix of scalar covariates

**Note**

The dataset has been preprocessed and standardized for modeling purposes. Time series lengths and grid structures are aligned for all stations.

**References**

Maranzano, M., and Algieri, C. (2024). ARPALData: An R package for retrieving and analyzing air quality and weather data from ARPA Lombardia (Italy). *Environmental and Ecological Statistics*, 31(2), 187-218.

**Examples**

```
## NOTE: This example involves ultra high-dimensional functional data.
## Running the model may require a PC with at least 64 GB of RAM.

# data(air_data)

# y <- air_data$y
```

```

# x <- air_data$x
# xscl <- air_data$xscl

# ntot <- dim(y)[1]
# ntrain <- 1000
# ntest <- ntot - ntrain

# train_ind <- sample(1:ntot, ntrain, replace = FALSE)

# Training sample
# y_train <- y[train_ind, ]
# y_test <- y[-train_ind, ]
# xscl_train <- xscl[train_ind, ]
# xscl_test <- xscl[-train_ind, ]

# Test sample
# x_train <- x_test <- vector("list", length = 5)
# for (j in 1:5) {
#   x_train[[j]] <- x[[j]][train_ind, ]
#   x_test[[j]] <- x[[j]][-train_ind, ]
# }

# Train HDTN approach
# nfof_model <- fonf_fit(resp = y_train, func_cov = x_train, scalar_cov = xscl_train)

# Obtain predictions with conformal prediction intervals
# band <- fonf_predict(nfof_model,
#                       func_cov_new = x_test,
#                       scalar_cov_new = xscl_test,
#                       interval = "conformal")

```

dgp\_mixed

*Simulate Mixed Functional–Scalar Data for Function–on–Function Regression*

## Description

Generates synthetic datasets that mimic the structure analysed in *Beyaztas, Bakicierler Sezer, Inan (2025)*: a functional response observed on a dense grid, multiple functional predictors, and multiple scalar predictors. Two regimes are available:

- "linear": response driven solely by linear integral effects.
- "nonlinear": adds strong, smooth nonlinear interactions (functional  $\times$  functional, functional  $\times$  scalar, scalar  $\times$  scalar) scaled to dominate the linear part.

Every dataset also includes Ornstein–Uhlenbeck (OU) process noise to emulate realistic autocorrelated measurement error.

## Usage

```

dgp_mixed(n, j,
          model = c("linear", "nonlinear"),
          n_func = 5L,
          n_scl = 3L,
          seed = NULL)

```

### Arguments

n	Number of subjects/curves.
j	Grid length $M$ ; all functional objects are evaluated on <code>seq(0, 1, length.out = j)</code> .
model	Character string choosing the data-generating mechanism; partial matching is supported.
n_func	Number $P$ of functional predictors to simulate.
n_scl	Number $d_z$ of scalar predictors.
seed	Optional integer for reproducibility ( <code>set.seed(seed)</code> ).

### Details

Let  $X_i^{(p)}(s)$  denote functional predictor  $p = 1, \dots, P$  and  $\mathbf{Z}_i \in \mathbb{R}^{d_z}$  the scalar predictor vector for subject  $i \in \{1, \dots, n\}$ . The latent *signal* generating the functional response on grid points  $t_1, \dots, t_M$  is

$$\eta_i(t) = \underbrace{\sum_{p=1}^P \int_0^1 X_i^{(p)}(s) \beta_p(s, t) ds}_{\text{linear functional effects}} + \underbrace{\mathbf{Z}_i^\top \gamma(t)}_{\text{scalar effects}} + \mathcal{NL}_i(t),$$

where  $P = 3$  “strong” functional predictors drive the linear component,  $\beta_p$  are pre-specified wavy B-spline surfaces (see code), and  $\gamma(t)$  are smooth one-dimensional bases. If `model == "linear"` we set  $\mathcal{NL}_i(t) \equiv 0$ ; otherwise  $\mathcal{NL}_i(t)$  equals the *sum* of five carefully designed nonlinear terms (quadratic functional interactions, sinusoidal transforms, scalar–functional products, etc.) rescaled so that

$$\text{SD}\{\mathcal{NL}\} \approx 2 \text{SD}\{\text{linear signal}\}.$$

The *observed* response is

$$Y_i(t_m) = \eta_i(t_m) + \varepsilon_i(t_m), \quad \varepsilon_i \sim \text{OU}(0, \alpha = 3, \sigma = 0.7),$$

i.e. an Ornstein–Uhlenbeck process discretised on the same grid. Noise is scaled so that its empirical standard deviation equals 10% of the signal standard deviation.

### Value

A named list:

y	$n \times j$ matrix of noisy responses $Y_i(t_m)$ .
yt	$n \times j$ matrix of true signals $\eta_i(t_m)$ .
x	list of length $P$ ; each element is an $n \times j$ matrix of functional predictors $X_i^{(p)}(s_m)$ .
x.scl	$n \times d_z$ numeric matrix of scalar predictors $\mathbf{Z}_i$ .
meta	List containing grids <code>sx</code> , <code>sy</code> , true beta surfaces ( <code>beta</code> ), and the <code>model</code> flag.

### Note

Requires the suggested package **goffda** for OU noise generation.

**Author(s)**

Ufuk Beyaztas, Gizel Bakicierler Sezer, Deniz Inan

**References**

Kokoszka, P. and Reimherr, M.~(2017). *Introduction to Functional Data Analysis*. Chapman and Hall/CRC.

**See Also**

[fonf\\_fit](#), [r\\_ou](#)

**Examples**

```
## Not run:
# -----
# Simulate a nonlinear training set and inspect its structure
# -----
# n <- 100                                # sample size
# simdata <- dgp_mixed(n, 101, model = "nonlinear")

# y    <- simdata$y                      # noisy response curves
# yt   <- simdata$yt                     # true (noise-free) curves
# x    <- simdata$x                      # list of functional predictors
# xscl <- simdata$x.scl                  # scalar predictors
## End(Not run)
```

---

fonf\_fit

*Fit a Hybrid Deep Tensor Network for Function-on-Function Regression with Mixed Predictors*

---

**Description**

Trains the hybrid deep tensor network (HDTN) of *Beyaztas, Bakicierler Sezer, Inan (2025)* to predict a *functional response* from multiple functional and scalar covariates. The model combines a first-layer tensor-product B-spline representation (capturing linear functional effects) with fully-connected dense layers (capturing higher-order nonlinear interactions) and supplies finite-sample, distribution-free prediction bands via conformal inference.

**Usage**

```
fonf_fit(resp,
  func_cov,
  scalar_cov      = NULL,
  nbasis_y        = NULL,
  nbasis_x        = NULL,
  hidden_layers   = 2,
  neurons_per_layer = c(32, 32),
  activations_in_layers = c("relu", "linear"),
  epochs          = 100,
  batch_size      = 32,
  val_split       = 0.1,
```

```

learning_rate      = 1e-3,
patience_param    = 15,
dropout_rate       = 0.1,
l2_lambda          = 1e-4,
cal_prop           = 0.2,
alpha              = 0.2,
verbose            = 1)

```

### Arguments

resp	$n \times M$ numeric matrix; row $i$ stores the functional response $Y_i(t_m)$ evaluated on a common grid $t_1, \dots, t_M \subset \mathcal{I}_y$ .
func_cov	List of length $P$ . Element $p$ is an $n \times G_p$ matrix that holds the functional predictor $X_i^{(p)}(s_{pj})$ on its own grid $s_{p1}, \dots, s_{pG_p} \subset \mathcal{I}_{x_p}$ .
scalar_cov	Optional $n \times d_z$ numeric matrix of scalar covariates $\mathbf{Z}_i \in \mathbb{R}^{d_z}$ .
nbasis_y	Number of B-spline basis functions for the response domain ( $K_y$ ); chosen automatically when NULL.
nbasis_x	Integer vector of length $P$ ; element $p$ gives the number of input-domain basis functions $K_x^{(p)}$ . Automatically selected when NULL.
hidden_layers	Number $R$ of fully-connected hidden layers.
neurons_per_layer	Vector of length <code>hidden_layers</code> giving the width ( $D_r$ ) of each dense layer.
activations_in_layers	Character vector of length <code>hidden_layers</code> with Keras activation names ("relu", "tanh", etc.).
epochs	Maximum training epochs.
batch_size	Mini-batch size for stochastic optimisation.
val_split	Proportion of the <i>training rows</i> (not subjects) held out for on-line validation during training.
learning_rate	Initial learning rate for the Adam optimiser (with cosine decay scheduler).
patience_param	Early-stopping patience; training stops when validation loss fails to improve for this many epochs.
dropout_rate	Dropout probability applied after every dense hidden layer.
l2_lambda	$\ell_2$ (ridge) penalty applied to dense-layer weights.
cal_prop	Proportion of subjects set aside for the conformal <i>calibration</i> set.
alpha	Mis-coverage level for conformal prediction bands (e.g. 0.2 yields 80% bands).
verbose	Passed to keras; larger values give more console output.

### Details

Let  $Y_i(t) \in L^2(\mathcal{I}_y)$  be the response curve for subject  $i$ ,  $X_i^{(p)}(s) \in L^2(\mathcal{I}_{x_p})$ ,  $p \in \{1, \dots, P\}$ , the functional predictors, and  $\mathbf{Z}_i \in \mathbb{R}^{d_z}$  scalar predictors. The HDTN targets the nonlinear FoFR model

$$Y_i(t) = g \left\{ \beta_0(t) + \sum_{p=1}^P \langle X_i^{(p)}, \beta_p(\cdot, t) \rangle_{L^2} + \mathbf{Z}_i^\top \theta(t) \right\} + \varepsilon_i(t), \quad t \in \mathcal{I}_y,$$

with identity link  $g(u) = u$  in the present implementation.

**Tensor-product layer.** Each bivariate coefficient surface is expanded

$$\beta_p(s, t) = \sum_{k=1}^{K_x^{(p)}} \sum_{\ell=1}^{K_y} w_{k\ell}^{(p)} \phi_k^{(p)}(s) \psi_\ell(t),$$

yielding first-layer weights  $w_{k\ell}^{(p)}$  to be learned. Subject-specific functional features are  $\tilde{\varphi}_{ik}^{(p)} = \langle X_i^{(p)}, \phi_k^{(p)} \rangle_{L^2}$  and  $u_{i\ell}^{(p)} = \sum_k \tilde{\varphi}_{ik}^{(p)} w_{k\ell}^{(p)}$ .

**Dense layers.** The concatenated feature vector  $\mathbf{h}_i^{(0)} = (u_{i\cdot}^{(1)\top}, \dots, u_{i\cdot}^{(P)\top}, \mathbf{Z}_i^\top)^\top$  is propagated through  $R$  fully-connected layers  $\mathbf{h}_i^{(r)} = \sigma_r(W^{(r)}\mathbf{h}_i^{(r-1)} + b^{(r)})$  with dropout and ridge penalty  $\lambda \|W^{(r)}\|_F^2/2$ . The final linear layer outputs  $\hat{\eta}_i(t_m)$ , giving  $\hat{Y}_i(t_m) = \hat{\eta}_i(t_m)$  for identity link.

**Loss and optimisation.** Training minimises mean integrated squared error plus the ridge penalty,  $\mathcal{L} = \frac{1}{nM} \sum_{i=1}^n \sum_{m=1}^M \{Y_i(t_m) - \hat{Y}_i(t_m)\}^2 + \frac{\lambda}{2} \sum_{r=1}^R \|W^{(r)}\|_F^2$ , via Adam with cosine-decay learning rate.

**Conformal prediction.** A calibration subset  $C$  (size `cal_prop * n` subjects) yields residuals  $e_{jm} = |Y_j(t_m) - \hat{Y}_j(t_m)|$ . The empirical  $1 - \alpha$  quantile  $q_{1-\alpha}$  of `length(C)*M` pooled residuals provides a *constant-width*  $(1 - \alpha)$  band  $[\hat{Y}_i(t_m) - q_{1-\alpha}, \hat{Y}_i(t_m) + q_{1-\alpha}]$  for every new subject  $i$  and grid point  $t_m$ . Finite-sample marginal coverage is guaranteed (Lei, G'Sell, *et al.*, 2018).

## Value

An object of class "fonf\_dl" to be consumed by [fonf\\_predict](#):

<code>model</code>	Trained Keras model.
<code>center, scale</code>	Vectors used to standardise the design matrix.
<code>py</code>	Number of response grid points ( $M$ ).
<code>nbasis_y, nbasis_x</code>	Basis dimensions actually used.
<code>q_hat</code>	Half-width $q_{1-\alpha}$ of the conformal prediction band.
<code>alpha</code>	User-supplied mis-coverage level.
<code>history</code>	<code>keras_training_history</code> object returned by <code>fit()</code> .

## Note

Requires TensorFlow/Keras (tested with TensorFlow  $\geq 2.16$ ).

## Author(s)

Ufuk Beyaztas, Gizel Bakicierler Sezer, Deniz Inan

## References

Lei, J., G'Sell, M., Rinaldo, A., Tibshirani, R. and Wasserman, L. (2018). *Distribution-Free Predictive Inference for Regression*. *Journal of the American Statistical Association*, 113(523), 1094-1111.

## See Also

[fonf\\_predict](#), [keras\\_model\\_sequential](#)

## Examples

```
## Not run:
# -----
# 1. Simulate training data
# -----
# n <- 100
# simdata <- dgp_mixed(n, 101, model = "nonlinear")

# y   <- simdata$y
# yt  <- simdata$yt      # (true curves, if needed)
# x   <- simdata$x       # list of functional predictors
# xscl <- simdata$x.scl  # scalar predictors

# -----
# 2. Fit the hybrid deep tensor network
# -----
# nfof_model <- fonf_fit(resp      = y,
#                        func_cov   = x,
#                        scalar_cov = xscl)
## End(Not run)
```

---

fonf\_fnc

---

*Extract and Visualise Estimated Coefficient Surfaces*


---

## Description

Given a fitted `fonf_fit` model, this helper recovers the bivariate coefficient surfaces  $\hat{\beta}_p(t, s)$  that link functional predictor  $X^{(p)}(s)$  to the functional response  $Y(t)$ , evaluates them on user-supplied grids, and—optionally—renders compact, printer-friendly 3-D perspective plots via **plot3D**.

## Usage

```
fonf_fnc(model,
          y_grid      = NULL,
          x_grid      = NULL,
          agg_fun      = mean,
          plot         = TRUE,
          grid_len     = 101,
          view_theta   = 40,
          view_phi     = 2,
          surface_col  = "royalblue",
          border_col   = "black",
          shade_fac    = 0.5,
          title_mgp    = c(2.8, 0.8, 0),
          ...)
```

## Arguments

<code>model</code>	Object of class "fonf_dl" returned by <code>fonf_fit</code> .
<code>y_grid</code>	Numeric vector of evaluation points $t_1, \dots, t_M$ in the interval $[0, 1]$ . Defaults to <code>seq(0, 1, length.out = grid_len)</code> .



<code>x_grid</code>	List of length $P$ ; element $p$ is the grid $s_{k1}, \dots, s_{kL_k}$ for predictor $p$ . Defaults to a <code>length_grid_len</code> equi-spaced grid for every predictor.
<code>agg_fun</code>	Function used to collapse the first-layer tensor weights into a single set of B-spline coefficients. With the default mean: $\hat{\beta}_p(t, s) = \frac{1}{D_1} \sum_{d=1}^{D_1} w_d^{(p)} B_y(t) B_x^{(p)}(s).$
<code>plot</code>	Logical; if TRUE (default) one 3-D plot is produced for each predictor.
<code>grid_len</code>	Length of the default equi-spaced grids when <code>y_grid</code> or an element of <code>x_grid</code> is NULL.
<code>view_theta, view_phi</code>	Viewing angles (in degrees) passed to <a href="#">persp3D</a> .
<code>surface_col, border_col, shade_fac</code>	Graphical parameters for the surfaces.
<code>title_mgp</code>	Margin settings for titles (argument passed to <code>par(mgp = ...)</code> inside each plot).
<code>...</code>	Currently ignored, reserved for future extensions.

### Details

Let  $B_y(t)$  be the vector of  $K_y$  B-spline basis functions for  $t$  and let  $B_x^{(p)}(s)$  be the vector of  $K_x^{(p)}$  basis functions for  $s$  in predictor  $p$ . The first layer of the network stores a weight matrix  $W^{(1)}$  whose rows correspond to the spline products  $B_y(t) B_x^{(p)}(s)$ . For each predictor  $p$  we:

1. Extract the contiguous block of  $K_y K_x^{(p)}$  weights.
2. Fold it into a  $K_y \times K_x^{(p)}$  matrix  $C_k$ .
3. Evaluate  $\hat{\beta}_p(t, s) = B_y(t)^\top C_k B_x^{(p)}(s)$  on the requested grids.

If `plot = TRUE` every surface is drawn with [persp3D](#); otherwise the numeric matrices are returned silently.

### Value

Invisibly returns a list with components:

<code>beta_hat</code>	List of length $P$ ; element $p$ is a <code>length(y_grid)</code> by <code>length(x_grid[[p]])</code> matrix containing $\hat{\beta}_p(t, s)$ .
<code>y_grid</code>	Evaluation grid for $t$ .
<code>x_grid</code>	Evaluation grids for $s$ .
<code>plots</code>	List of closure functions; calling <code>plots[[p]]()</code> re-draws the $p$ -th surface.

### Note

Requires the suggested packages **fda**, **keras**, and **plot3D**. An error is thrown if any are missing.

### Author(s)

Ufuk Beyaztas, Gizel Bakicierler Sezer, Deniz Inan

### See Also

[fonf\\_fit](#), [persp3D](#), [eval.basis](#)

**Examples**

```
## Not run:
# -----
# Fit a small model and visualise its coefficient surfaces
# -----
# simdata <- dgp_mixed(100, 101, model = "nonlinear")
#
# y <- simdata$y
# x <- simdata$x
# xscl <- simdata$x.scl
#
# mdl <- fonf_fit(resp = y, func_cov = x, scalar_cov = xscl)
#
# Plot the 2-D functional weights
# surf_obj <- fonf_fnc(model = mdl)
## End(Not run)
```

fonf\_param\_grid

*Create a Parameter Grid Without Factors***Description**

Generates a data frame from all combinations of input vectors, ensuring that character variables are not converted to factors. This is a convenient wrapper around `expand.grid(..., stringsAsFactors = FALSE)`.

**Usage**

```
fonf_param_grid(...)
```

**Arguments**

... One or more vectors, factors, or lists to be combined into a data frame of parameter combinations.

**Details**

This function simplifies the creation of parameter grids for tuning models. Unlike the base R `expand.grid`, it ensures that character vectors remain as characters, which is often desirable when building machine learning or neural network models where hyperparameter names or tags are string-based.

**Value**

A data frame containing one row for each combination of the supplied vectors.

**Note**

This function is mainly used to build cross-product grids of model parameters when performing grid search for tuning.

**Author(s)**

Ufuk Beyaztas, Gizel Bakicierler Sezer, Deniz Inan

fonf\_predict

*Predict With a Trained Hybrid Deep Tensor Network***Description**

Generates point predictions and, optionally, distribution-free conformal prediction bands for a [fonf\\_fit](#) model given new functional and scalar covariates.

**Usage**

```
fonf_predict(object,
             func_cov_new,
             scalar_cov_new = NULL,
             interval = c("none", "conformal"))
```

**Arguments**

object	An object of class "fonf_dl" returned by <a href="#">fonf_fit</a> .
func_cov_new	List of length $P$ ; element $p$ is an $n_{\text{new}} \times G_p$ matrix that stores the new subjects' functional predictor $X_{\text{new},i}^{(p)}(s_{pj})$ .
scalar_cov_new	Optional $n_{\text{new}} \times d_z$ numeric matrix of scalar covariates $Z_{\text{new},i}$ . If the original model was fitted without scalar covariates, set this to NULL.
interval	Type of prediction to return: "none" Point estimates only. "conformal" Point estimates plus lower/upper bands that achieve marginal coverage $1 - \alpha$ (see Details).

**Details**

Let  $f_{\text{hat}}$  denote the trained network and let  $(x_i, z_i)$  be the design-matrix rows for a new subject, built with the same tensor-product bases used in training. `fonf_predict`:

1. Builds the design matrix with `design_matrix_build`.
2. Standardises it with the training means and scales stored in the fitted object.
3. Computes the predictions, then reshapes the vector into an  $n_{\text{new}} \times M$  matrix, where  $M = \text{object\$py}$ .

**Conformal bands:** If `interval == "conformal"` the returned bands are  $\hat{Y}_i(t_m) \pm q_{1-\alpha}$ , where  $q_{1-\alpha}$  is the  $(1 - \alpha)$ -quantile of the absolute residuals on the calibration set chosen during training. Split-conformal theory (Lei *et al.*, 2018) guarantees marginal coverage  $1 - \alpha$  at each grid point, without distributional assumptions beyond independent rows.

**Value**

`interval = "none"` A numeric matrix with `nrow = n_new` and `ncol = py` containing point predictions.

`interval = "conformal"` A list with components `mean`, `lower`, and `upper`, each a matrix whose columns are named "t1", "t2", ..., "tM".

**Note**

The feature grids supplied here must match those used at training time, and the length of the prediction vector must be divisible by `object$py`; otherwise an error is raised.

**Author(s)**

Ufuk Beyaztas, Gizel Bakicierler Sezer, Deniz Inan

**References**

Lei, J., G'Sell, M., Rinaldo, A., Tibshirani, R. & Wasserman, L. (2018). Distribution-Free Predictive Inference for Regression. *Journal of the American Statistical Association*, 113(523), 1094–1111.

**See Also**

[fonf\\_fit](#), [dgp\\_mixed](#)

**Examples**

```
## Not run:
# -----
# 1. Train a model on simulated data
# -----
# simdata <- dgp_mixed(100, 101, model = "nonlinear")
# mdl <- fonf_fit(resp = simdata$y,
#               func_cov = simdata$x,
#               scalar_cov = simdata$x.scl)
#
# -----
# 2. Predict on new subjects with conformal bands
# -----
# test <- dgp_mixed(250, 101, model = "nonlinear")
# band <- fonf_predict(mdl,
#                   func_cov_new = test$x,
#                   scalar_cov_new = test$x.scl,
#                   interval = "conformal")
# pred_mean <- band$mean
## End(Not run)
```

---

fonf\_tune

*Simple Grid-Search Hyper-parameter Tuner for fonf\_fit*

---

**Description**

Evaluates every row of a user-supplied hyper-parameter grid via  $K$ -fold cross-validation (`fonf_cv`), reports the cross-validated mean-squared error (CV-MSE), and then refits the best combination on the *full* data set. The implementation is deliberately lightweight—single-core, base R only—so it runs on any system where **fonf** installs.

**Usage**

```
fonf_tune(grid,
          resp, func_cov, scalar_cov = NULL,
          nfolds = 5)
```

## Arguments

grid	A data frame created by <code>fonf_param_grid</code> describing the hyper-parameter combinations to be tested. Each column corresponds to an argument of <code>fonf_fit</code> .
resp	Numeric matrix of size $n \times p_y$ ; functional response curves sampled on a common grid.
func_cov	List of length $P$ ; element $p$ is an $n \times G_p$ matrix containing the $p$ -th functional predictor.
scalar_cov	Optional $n \times q$ numeric matrix of scalar predictors. Omit or set to NULL if none.
nfolds	Number of folds used by <code>fonf_cv</code> (default 5).

## Details

**How it works:** For each row  $g = 1, \dots, G$  of `grid`:

1. The row is coerced to a named list of arguments compatible with `fonf_fit` (vectors such as `neurons_per_layer` are replicated to the correct length).
2. `fonf_cv` performs  $K$ -fold CV and stores the average test MSE in `cv_vec[g]`.

After all rows are processed the index of the minimum  $CV-MSE$  is selected, the corresponding parameter list is re-sanitised by `sanitize_basis()`, and a final model is fitted on the whole data via `fonf_fit` with `verbose = TRUE`.

## Value

**results** Data frame combining the original `grid` and a new column `CV_MSE` (lower is better).

**best\_params** Named list containing the best-performing hyper-parameters in a format ready for `fonf_fit`.

**best\_model** Object of class "fonf\_d1" fitted on the full data with `best_params`.

## Note

- This function is single-threaded. For large grids consider a parallel wrapper (e.g. `future.apply`).
- Internal helpers `flatten1()` and `sanitize_basis()` are not exported but are documented in the source code.

## Author(s)

Ufuk Beyaztas, Gizel Bakicierler Sezer, Deniz Inan

## Examples

```
## Not run:
# -----
# 1. Simulate training data
# -----
# n <- 100
# simdata <- dgp_mixed(n, 101, model = "nonlinear")
#
# y <- simdata$y
# x <- simdata$x
# xscl <- simdata$x.scl
#
# -----
```

```

# 2. Construct a small hyper-parameter grid
# -----
# grid <- fonf_param_grid(
#   hidden_layers      = c(1, 2),
#   neurons_per_layer  = list(32, c(64, 32)),
#   activations_in_layers = list("relu", c("relu", "linear")),
#   learning_rate       = c(1e-3, 5e-4),
#   epochs              = 25
# )
# -----
# 3. Tune and refit
# -----
# tune_out <- fonf_tune(grid,
#   resp      = y,
#   func_cov  = x,
#   scalar_cov = xscl,
#   nfolds    = 5)
#
# tune_out$results      # CV table
# best_mod <- tune_out$best_model
## End(Not run)

```

# Index

`air_data`, [1](#)

`dgp_mixed`, [3](#), [12](#)

`eval.basis`, [9](#)

`fonf_fit`, [5](#), [5](#), [8](#), [9](#), [11–13](#)

`fonf_fnc`, [8](#)

`fonf_param_grid`, [10](#)

`fonf_predict`, [7](#), [11](#)

`fonf_tune`, [12](#)

`keras_model_sequential`, [7](#)

`persp3D`, [9](#)

`r_ou`, [5](#)